

# Twitter Sentiment Analysis

## Problem Statement

The task is to build a model that will determine the tone (neutral, positive, negative) of the text. To do this, you will need to train the model on the existing data (train.csv). The resulting model will have to determine the class (neutral, positive, negative) of new texts (test data that were not used to build the model) with maximum accuracy.

## Performance metric

Source: <https://www.kaggle.com/competitions/twitter-sentiment-analysis2/overview/evaluation>  
(<https://www.kaggle.com/competitions/twitter-sentiment-analysis2/overview/evaluation>)

Metric: F1-score

Loading the required libraries

```
In [2]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from wordcloud import WordCloud
import nltk
from nltk.corpus import stopwords
import warnings
warnings.filterwarnings("ignore")
from sklearn.feature_extraction.text import CountVectorizer
from nltk.stem.snowball import SnowballStemmer
import re
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
import pickle
from tqdm import tqdm
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
import seaborn as sns
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import LinearSVC
from tabulate import tabulate
import tensorflow as tf
from tensorflow.keras.layers import Dense, Input, Activation
from tensorflow.keras.models import Model
import random as rn
from sklearn.metrics import f1_score, roc_auc_score
from tensorflow.keras.callbacks import LearningRateScheduler
from tensorflow.keras.callbacks import EarlyStopping
import datetime
import os
import tensorflow_addons as tfa
from tensorflow.keras import backend as K
import tensorflow_hub as hub
import tensorflow_text as text
```

Data source: <https://www.kaggle.com/competitions/twitter-sentiment-analysis2/data>  
(<https://www.kaggle.com/competitions/twitter-sentiment-analysis2/data>)

## Data description:

### File descriptions

train.csv - the training set

test.csv - the test set

### Data fields:

ItemID - id of tweet

SentimentText - text of the tweet

Sentiment - sentiment

0 - negative

1 - positive

## Loading the data

```
In [3]: train = pd.read_csv(r"D:\Twitter Sentiment Analysis\Dataset\train.csv", encoding='latin-1')
train.head(4)
```

Out[3]:

	ItemID	Sentiment	SentimentText
0	1	0	is so sad for my APL frie...
1	2	0	I missed the New Moon trail...
2	3	1	omg its already 7:30 :O
3	4	0	.. Omgaga. Im sooo im gunna CRy. I'...

```
In [4]: print(train.shape)

(99989, 3)
```


## Data Visualization

```
In [5]: print("Null data in the train dataset:\n", train.isnull().any())
```

```
Null data in the train dataset:  
ItemID          False  
Sentiment        False  
SentimentText     False  
dtype: bool
```

```
In [6]: train['Sentiment'].value_counts().plot.bar(color='#e8948e', figsize=(6,4))
```


```
Out[6]: <AxesSubplot:>
```



Since, the dataset is slightly imbalanced and is large enough resulting in the increase of training time for deep learning models we will use random sampling to make both class labels equal

```
In [18]: train['SentimentText'].str.len().plot.hist(color='#E57F84', figsize=(6,4))
```

```
Out[18]: <AxesSubplot:ylabel='Frequency'>
```



```
In [19]: nltk.download('stopwords')
stop_words = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\anike\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
In [20]: stop_words.remove('not')
stop_words.remove('no')
print("Stop Words :\n", stop_words)
```

Stop Words :

```
{'at', 'will', 'am', 'has', 'aren', 'if', 'as', 'off', "shouldn't", 'yourself', 'or', 'than', 'me', 'your', 'out', 'mightn', 'did', 'had', 'these', 'and', 'to', 'down', 'd', "mightn't", 'is', 'do', 'after', 'again', 'on', 'above', 'more', 'what', "don't", 'his', 'where', 'wasn', "shan't", 's', 'for', 're', 'yourselves', "won't", "haven't", 'between', 'from', "couldn't", 'ma', 'our', 'through', "should've", 'until', 'her', "weren't", 'she', 'mustn', "wasn't", 've', 'doesn', 'being', 'very', 'hasn', 'once', 'why', 'herself', 'just', 'it', 'itself', 'm', 'he', 'by', 'didn', 'shan', 'don', "that'll", "you'd", 'does', "didn't", 'isn', 'before', 'weren', 'during', 'been', 'other', "aren't", 'how', 'any', 'each', 'was', 'too', 'i', "it's", "hasn't", 'now', 'they', 'himself', 'an', 'themselves', "you've", 'then', 'some', 'y', 'here', 'theirs', 'myself', 'those', 'o', "isn't", 'against', "wouldn't", 'own', 'wouldn', 'yours', 'over', 'under', 'hers', 'most', "you're", 'whom', 'about', "mustn't", 'nor', 'such', 'll', 'ain', 'few', 'haven', 'them', 'be', 'were', 'a', 't', 'of', 'won', 'all', 'that', 'both', 'their', 'have', 'we', 'hadn', 'needn', "you'll", 'you', 'its', 'below', 'while', "doesn't", 'should', 'having', 'the', 'couldn', "she's", 'further', 'there', 'so', 'same', 'up', 'in', 'because', 'with', 'my', 'him', 'which', 'doing', 'ours', 'ourselves', 'only', "needn't", 'but', 'when', "hadn't", 'shouldn', 'into', 'this', 'are', 'who', 'can'}
```

```
In [21]: cv = CountVectorizer(stop_words='english')
words = cv.fit_transform(train['SentimentText'])

sum_words = words.sum(axis=0)

word_freq = [(word, sum_words[0,i]) for word,i in cv.vocabulary_.items()]
word_freq = sorted(word_freq, key=lambda x:x[1], reverse=True)
```



```
In [23]: positive_words = ' '.join([word for word in train['SentimentText'][train['Sentiment']==1]])
positive_wc = WordCloud(background_color='White', width = 1000, height = 1000)
positive_wc.generate(positive_words)
plt.figure(figsize=(8, 8))
plt.axis('off')
plt.imshow(positive_wc)
plt.title("Word Cloud for positive words", fontsize=40)
```

Out[23]: Text(0.5, 1.0, 'Word Cloud for positive words')

## Word Cloud for positive words







```
In [27]: def decontracted(phrase):
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)
    phrase = re.sub(r"\ve", " have", phrase)
    phrase = re.sub(r"\m", " am", phrase)
    return phrase
```

```
In [28]: def text_preprocessing(text):
    text = re.sub('@[\s]+', ' ', text)
    text = re.sub('((www\.[^\s]+)|(https?://[^\s]+))', ' ', text)
    text = re.sub('[^a-zA-Z\n]', ' ', text)
    text = re.sub('\s+', ' ', text)
    text = re.sub("\n", " ", text)
    text = re.sub("\t", " ", text)
    text = re.sub(",", " ", text)
    #text = ' '.join(words for words in text.split(" ") if len(words)>2)
    text = ' '.join(words for words in text.split(" "))
    text = decontracted(text)
    text = re.sub(r'(\.)\1{3,}', r'\1', text)
    text = re.sub(r'[\W\s]', ' ', text)
    return text.strip().lower()
```

```
In [29]: def stop_and_stem(text):
    text = ' '.join(words.lower() for words in text.split(" ") if words not in
stop_words)
    word_list = nltk.word_tokenize(text)
    text = ' '.join(snow_stemmer.stem(word) for word in word_list)
    text = ' '.join(words for words in text.split(" ") if len(words)>2)
    return text
```

```
In [30]: def stop_and_lemmatize(text):
    text = ' '.join(words.lower() for words in text.split(" ") if words not in
stop_words)
    word_list = nltk.word_tokenize(text)
    text = ' '.join(lemmatizer.lemmatize(word) for word in word_list)
    text = ' '.join(words for words in text.split(" ") if len(words)>2)
    return text
```

```
In [106]: def gt_2(text):
    text = ' '.join(words for words in text.split(" ") if len(words)>2)
    return text
```

```
In [31]: train["clean_SentimentText"] = train["SentimentText"].apply(text_preprocessing
)
```



```
In [37]: Y = train[['Sentiment']]
X = train[['lemmatized_SentimentText']]
```

```
In [38]: print(X.shape)
print(Y.shape)
```

```
(99989, 1)
(99989, 1)
```

```
In [39]: X.head(2)
```

Out[39]:

	lemmatized_SentimentText
0	sad apl friend
1	missed new moon trailer

```
In [40]: Y.head(2)
```

Out[40]:

	Sentiment
0	0
1	0

## Splitting the data

```
In [41]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.33, stratify=Y)
```

```
In [42]: print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(66992, 1)
(66992, 1)
(32997, 1)
(32997, 1)
```

```
In [43]: X_train.head(2)
```

Out[43]:

	lemmatized_SentimentText
48200	mouthng said stupid people like think one gre...
19182	quot everyone become great quot one still isnt...

```
In [44]: X_test.head(2)
```

```
Out[44]:
```

	lemmatized_SentimentText
84535	reply
28022	thank much look resource

```
In [45]: tfidf_vectorizer = TfidfVectorizer(min_df=8, ngram_range=(1,3))
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train['lemmatized_SentimentText'].values)
X_test_tfidf = tfidf_vectorizer.transform(X_test['lemmatized_SentimentText'].values)
```

```
In [46]: print(X_train_tfidf.shape)
print(X_test_tfidf.shape)
```

```
(66992, 7590)
```

```
(32997, 7590)
```

```
In [47]: with open(r'D:\Twitter Sentiment Analysis\Dataset\glove_vectors', 'rb') as f:
        model = pickle.load(f)
        glove_words = set(model.keys())
```

```
In [48]: dictionary = dict(zip(tfidf_vectorizer.get_feature_names(), list(tfidf_vectorizer.idf_)))
tfidf_words = set(tfidf_vectorizer.get_feature_names())
```





```
In [52]: parameters = {"C": [10**x for x in range(-4,5)]}

model = LogisticRegression(penalty='l2', random_state=0)
clf_log = GridSearchCV(model, parameters, cv=10, n_jobs=-1)
clf_log.fit(X_train_tfidf, Y_train)

print("Best estimator=", clf_log.best_estimator_)
print("Best score =", clf_log.best_score_)
```

```
Best estimator= LogisticRegression(C=1, random_state=0)
Best score = 0.7501193673370719
```

```
In [53]: model_log = LogisticRegression(C= clf_log.best_params_['C'], penalty='l2', ran
dom_state=0)
model_log.fit(X_train_tfidf, Y_train)


Y_pred = model_log.predict(X_test_tfidf)

f1_score_log = f1_score(Y_test, Y_pred)
print("The f1 score for logistice regression model is ", f1_score_log)

cm_log = confusion_matrix(Y_test, Y_pred)
print("The confusion matrix on the test dataset for logistic regression:")
#print(cm_log)
sns.heatmap(cm_log, annot=True, fmt='d', cmap='Blues')
```

```
The f1 score for logistice regression model is 0.7869945383214954
The confusion matrix on the test dataset for logistic regression:
```

Out[53]: <AxesSubplot:>



## Naive Bayes

```
In [54]: multiNB = MultinomialNB()
parameters = {'alpha':[10 ** x for x in range(-3, 3)]}
clf_NB = GridSearchCV(multiNB, parameters, cv=10, n_jobs=-1)
clf_NB.fit(X_train_tfidf, Y_train)

print("Best estimator=", clf_NB.best_estimator_)
print("Best score =", clf_NB.best_score_)
```

```
Best estimator= MultinomialNB(alpha=1)
Best score = 0.7439543638725317
```

```
In [55]: model_NB = MultinomialNB(alpha= clf_NB.best_params_['alpha'])
model_NB.fit(X_train_tfidf, Y_train)


Y_pred = model_NB.predict(X_test_tfidf)

f1_score_NB = f1_score(Y_test, Y_pred)
print("The f1 score for Naive Bayes model is ", f1_score_NB)

cm_NB = confusion_matrix(Y_test, Y_pred)
print("The confusion matrix on the test dataset for Naive Bayes:")
#print(cm_NB)
sns.heatmap(cm_NB, annot=True, fmt='d', cmap='Blues')
```

```
The f1 score for Naive Bayes model is 0.7844408427876822
The confusion matrix on the test dataset for Naive Bayes:
```

Out[55]: <AxesSubplot:>



## Random Forest



```
In [56]: model_RF = RandomForestClassifier()
parameters={'min_samples_split':[10,12,14], 'n_estimators':[100, 200, 350]}
clf_RF = GridSearchCV(model_RF, parameters, cv=10, n_jobs=-1)
clf_RF.fit(X_train_tfidf, Y_train)

print("Best estimator=", clf_RF.best_estimator_)
print("Best score=", clf_RF.best_score_)
```

```
Best estimator= RandomForestClassifier(min_samples_split=10, n_estimators=350)
Best score= 0.7357594272257164
```

```
In [57]: model_RF = RandomForestClassifier(min_samples_split= clf_RF.best_params_['min_
samples_split'], n_estimators= clf_RF.best_params_['n_estimators'])
model_RF.fit(X_train_tfidf, Y_train)


Y_pred = model_RF.predict(X_test_tfidf)

f1_score_RF = f1_score(Y_test, Y_pred)
print("The f1 score for Random Forest model is ", f1_score_RF)

cm_RF = confusion_matrix(Y_test, Y_pred)
print("The confusion matrix on the test dataset for Random Forest:")
#print(cm_RF)
sns.heatmap(cm_RF, annot=True, fmt='d', cmap='Blues')
```

```
The f1 score for Random Forest model is 0.7770627919045148
The confusion matrix on the test dataset for Random Forest:
```

Out[57]: <AxesSubplot:>



## Gradient Boosting Decision Trees

```
In [58]: model_GBDT = GradientBoostingClassifier()
parameters = {'learning_rate':[0.1, 0.3, 0.4], 'n_estimators':[200, 300, 500]}
clf_GBDT = GridSearchCV(model_GBDT, parameters, cv=10, n_jobs=-1)
clf_GBDT.fit(X_train_tfidf, Y_train)

print("Best estimator=", clf_GBDT.best_estimator_)
print("Best score=", clf_GBDT.best_score_)
```

```
Best estimator= GradientBoostingClassifier(learning_rate=0.3, n_estimators=500)
Best score= 0.7374164555636507
```

```
In [59]: model_GBDT = GradientBoostingClassifier(learning_rate= clf_GBDT.best_params_['learning_rate'],
n_estimators= clf_GBDT.best_params_['n_estimators'])
model_GBDT.fit(X_train_tfidf, Y_train)


Y_pred = model_GBDT.predict(X_test_tfidf)

f1_score_GBDT = f1_score(Y_test, Y_pred)
print("The f1 score for Gradient Boosting Decision Tree model is ", f1_score_GBDT)

cm_GBDT = confusion_matrix(Y_test, Y_pred)
print("The confusion matrix on the test dataset for Gradient Boosting Decision Tree:")
#print(cm_GBDT)
sns.heatmap(cm_GBDT, annot=True, fmt='d', cmap='Blues')
```

```
The f1 score for Gradient Boosting Decision Tree model is 0.7870674537505866
The confusion matrix on the test dataset for Gradient Boosting Decision Tree:
```

Out[59]: <AxesSubplot:>



## Linear SVC

```
In [60]: model_SVC = LinearSVC()
parameters = {"C":[0.001, 0.01, 0.1, 1, 10]}
clf_SVC = GridSearchCV(model_SVC, parameters, cv=10, n_jobs=-1)
clf_SVC.fit(X_train_tfidf, Y_train)

print("Best estimator= ", clf_SVC.best_estimator_)
print("Best score= ", clf_SVC.best_score_)
```

```
Best estimator=  LinearSVC(C=0.1)
Best score=  0.7502835709495514
```

```
In [61]: model_SVC = LinearSVC(C=clf_SVC.best_params_["C"])
model_SVC.fit(X_train_tfidf, Y_train)


Y_pred = model_SVC.predict(X_test_tfidf)

f1_score_SVC = f1_score(Y_test, Y_pred)
print("The f1 score for Gradient linear SVC model is ", f1_score_SVC)

cm_SVC = confusion_matrix(Y_test, Y_pred)
print("The confusion matrix on the test dataset for linear SVC:")
#print(cm_SVC)
sns.heatmap(cm_SVC, annot=True, fmt='d', cmap='Blues')
```

```
The f1 score for Gradient linear SVC model is  0.7886519032554343
The confusion matrix on the test dataset for linear SVC:
```

```
Out[61]: <AxesSubplot:>
```



## Deep Learning Models

### Data preprocessing

```
In [109]: #Y = train[['Sentiment']]
#X = train[['clean_SentimentText']]
X = train[['preprocessed_SentimentText']]
```

```
In [110]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.20, stratify=Y)
X_train, X_cv, Y_train, Y_cv = train_test_split(X_train,Y_train, test_size=0.20, stratify=Y_train)
```

```
In [111]: print(X_train.shape, Y_train.shape)
print(X_cv.shape, Y_cv.shape)
print(X_test.shape, Y_test.shape)
```

```
(63992, 1) (63992, 1)
(15999, 1) (15999, 1)
(19998, 1) (19998, 1)
```

```
In [113]: tokenizer = tf.keras.preprocessing.text.Tokenizer(filters='!"#$%&()*+,-/:;<=>?@[\]\^`{|}~\t\n')
tokenizer.fit_on_texts(X_train['preprocessed_SentimentText'])

X_train_seq = tokenizer.texts_to_sequences(X_train['preprocessed_SentimentText'])
X_cv_seq = tokenizer.texts_to_sequences(X_cv['preprocessed_SentimentText'])
X_test_seq = tokenizer.texts_to_sequences(X_test['preprocessed_SentimentText'])

word_index = tokenizer.word_index
result = [len(x.split()) for x in X_train['preprocessed_SentimentText']]

print("Text informations:")
print("Max length: {}".format(np.max(result)))
print("Min length: {}".format(np.min(result)))
print("Mean length: {}".format(np.mean(result)))
```

```
Text informations:
Max length: 84
Min length: 0
Mean length: 9.180772596574572
```

```
In [114]: max_length= int(np.max(result))
print(max_length)
```

```
84
```

```
In [115]: Q3_len = int(X_train['preprocessed_SentimentText'].str.split().apply(len).describe()['75%'])
X_train['preprocessed_SentimentText'].str.split().apply(len).describe()
```

```
Out[115]: count      63992.000000
mean           9.180773
std            5.233915
min            0.000000
25%            5.000000
50%            8.000000
75%           13.000000
max           84.000000
Name: preprocessed_SentimentText, dtype: float64
```

```
In [116]: X_train_pad = tf.keras.preprocessing.sequence.pad_sequences(X_train_seq, maxlen=max_length)
X_cv_pad = tf.keras.preprocessing.sequence.pad_sequences(X_cv_seq, maxlen=max_length)
X_test_pad = tf.keras.preprocessing.sequence.pad_sequences(X_test_seq, maxlen=max_length)
```

```
In [117]: print(X_train_pad.shape)
print(X_cv_pad.shape)
print(X_test_pad.shape)
```

```
(63992, 84)
(15999, 84)
(19998, 84)
```

```
In [120]: %%time

#Reference from https://blog.keras.io/using-pre-trained-word-embeddings-in-a-keras-model.html
glove_file = r'D:\Twitter Sentiment Analysis\Dataset\glove.840B.300d.pkl'
glove = np.load(glove_file, allow_pickle=True)
```

```
EMBEDDING_DIM=300
embedding_matrix = np.zeros((len(word_index) + 1, EMBEDDING_DIM))
for word, i in word_index.items():
    embedding_vector = glove.get(word)
    if embedding_vector is not None:
        # words not found in embedding index will be all-zeros.
        embedding_matrix[i] = embedding_vector
print(embedding_matrix.shape)
```

```
(36910, 300)
Wall time: 7.76 s
```

```
In [121]: def Recall(y_true, y_pred):
    y_true = K.ones_like(y_true)
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    all_positives = K.sum(K.round(K.clip(y_true, 0, 1)))

    recall = true_positives / (all_positives + K.epsilon())
    return recall

def Precision(y_true, y_pred):
    y_true = K.ones_like(y_true)
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))

    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def F1_score(y_true, y_pred):
    precision = Precision(y_true, y_pred)
    recall = Recall(y_true, y_pred)
    return 2*((precision*recall)/(precision+recall+K.epsilon()))
```

## Sequential model using RNN with pre-trained word embeddings

```
In [122]: tf.random.set_seed(0)
np.random.seed(0)
```

```
In [124]: tf.compat.v1.keras.backend.get_session()
```

```
Out[124]: <tensorflow.python.client.session.Session at 0x2325a2a7c18>
```

```
In [125]: def build_rnn():
    model = tf.keras.models.Sequential()
    model.add(tf.keras.layers.Input(shape=max_length))
    model.add(tf.keras.layers.Embedding(len(tokenizer.word_index)+1, EMBEDDING_DIM, weights=[embedding_matrix], input_length=max_length, trainable=True))
    #model.add(tf.keras.layers.SimpleRNN(32, return_sequences=True, recurrent_dropout=0.5, bias_regularizer='L2'))
    model.add(tf.keras.layers.SimpleRNN(512, return_sequences=True, recurrent_dropout=0.3))
    model.add(tf.keras.layers.SimpleRNN(512, recurrent_dropout=0.3))
    model.add(tf.keras.layers.Dense(len(Y_train.columns), activation='sigmoid'))

    model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])
    return model
```

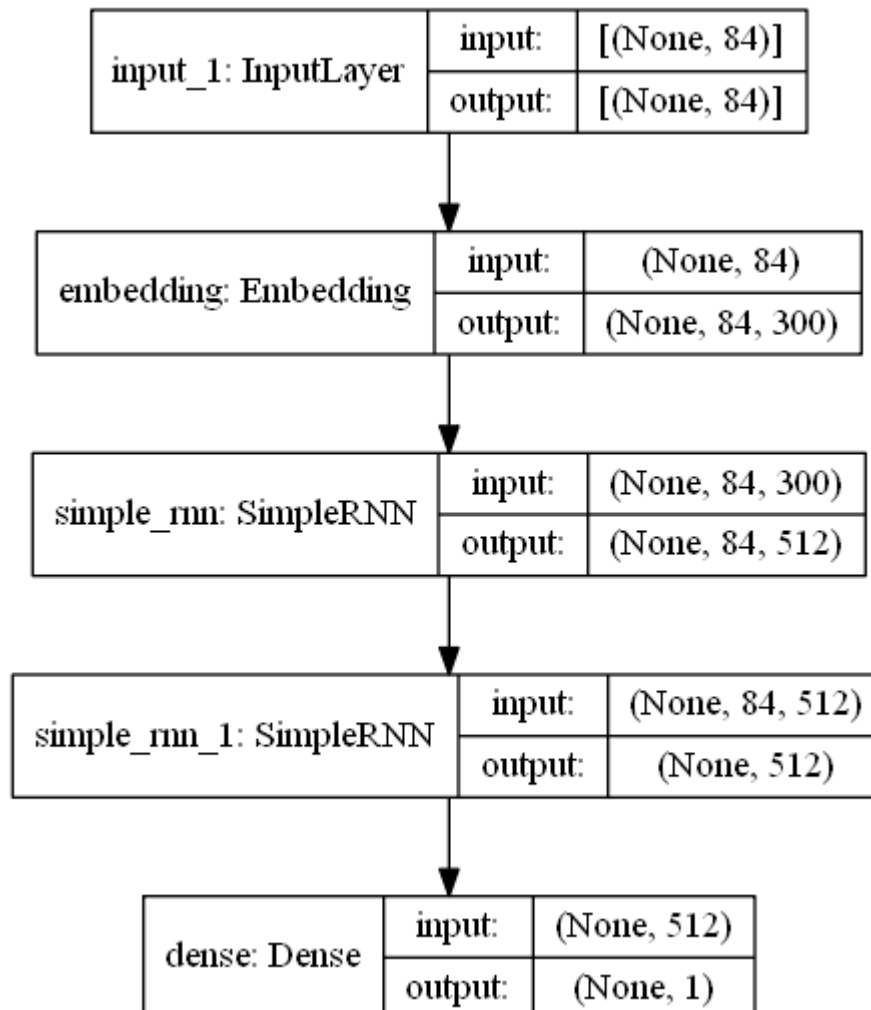
```
In [126]: model_rnn = build_rnn()
model_rnn.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 84, 300)	11073000
-----		
simple_rnn (SimpleRNN)	(None, 84, 512)	416256
-----		
simple_rnn_1 (SimpleRNN)	(None, 512)	524800
-----		
dense (Dense)	(None, 1)	513
=====		
Total params: 12,014,569		
Trainable params: 12,014,569		
Non-trainable params: 0		

```
In [127]: tf.keras.utils.plot_model(model_rnn,to_file='modelRNN.png', show_shapes=True,
show_layer_names=True)
```

Out[127]:



```
In [128]: checkpoint_path = "D:\Twitter Sentiment Analysis"
checkpoint_dir = os.path.dirname(checkpoint_path)
file_path=os.path.join(checkpoint_dir, "Twitter Sentiment Analysis", "checkpoi
nt", "weights-{epoch:02d}-{val_F1_score:.4f}")

earlystop = EarlyStopping(monitor='val_loss', patience=2, verbose=1)

model_checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(filepath=file_p
ath, save_weights_only=True, monitor='val_F1_score', mode='max', save_best_onl
y=True)

log_dir = os.path.join(checkpoint_dir, "Twitter Sentiment Analysis", "logs",'f
it', datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogra
m_freq=1, write_graph=True)
```

```
In [129]: %time
rnn_hist = model_rnn.fit(X_train_pad, Y_train, epochs=10, validation_data=(X_c
v_pad, Y_cv), batch_size=64,
                    callbacks=[earlystop, model_checkpoint_callback, tens
orboard_callback])
```

```
Epoch 1/10
1000/1000 [=====] - 729s 728ms/step - loss: 0.5686 -
accuracy: 0.7041 - F1_score: 0.7247 - Precision: 1.0000 - Recall: 0.5849 - va
l_loss: 0.5175 - val_accuracy: 0.7544 - val_F1_score: 0.7610 - val_Precision:
1.0000 - val_Recall: 0.6162
Epoch 2/10
1000/1000 [=====] - 798s 798ms/step - loss: 0.5032 -
accuracy: 0.7510 - F1_score: 0.7283 - Precision: 0.9990 - Recall: 0.5893 - va
l_loss: 0.5339 - val_accuracy: 0.7455 - val_F1_score: 0.7086 - val_Precision:
1.0000 - val_Recall: 0.5505
Epoch 3/10
1000/1000 [=====] - 807s 807ms/step - loss: 0.5137 -
accuracy: 0.7519 - F1_score: 0.7338 - Precision: 0.9970 - Recall: 0.5961 - va
l_loss: 0.6212 - val_accuracy: 0.7055 - val_F1_score: 0.5328 - val_Precision:
1.0000 - val_Recall: 0.3653
Epoch 00003: early stopping
Wall time: 38min 55s
```

```
In [130]: for metric, values in rnn_hist.history.items():
           print(metric, ":", values)
```

```
loss : [0.5685679912567139, 0.5031551718711853, 0.5136750936508179]
accuracy : [0.7040567398071289, 0.750968873500824, 0.7518908381462097]
F1_score : [0.7246702313423157, 0.7283213138580322, 0.7338063716888428]
Precision : [1.0, 0.9990000128746033, 0.996999979019165]
Recall : [0.5849397778511047, 0.5893124938011169, 0.5961004495620728]
val_loss : [0.5175146460533142, 0.5338994860649109, 0.6211705803871155]
val_accuracy : [0.754422128200531, 0.7455465793609619, 0.7054815888404846]
val_F1_score : [0.7609983682632446, 0.7085798382759094, 0.5328055620193481]
val_Precision : [1.0, 1.0, 1.0]
val_Recall : [0.6161696314811707, 0.5505436658859253, 0.36533334851264954]
```



```
In [131]: loss_rnn, accuracy_rnn, f1_score_rnn, precision_rnn, recall_rnn = model_rnn.evaluate(X_test_pad, Y_test, verbose=0)
print('F1 score for RNN model: %f' % f1_score_rnn)

F1 score for RNN model: 0.513728
```

```
In [132]: %load_ext tensorboard
```

```
In [133]: # Clear any logs from previous runs
!rm -rf ./logs/
```

'rm' is not recognized as an internal or external command,  
operable program or batch file.

```
In [134]: %tensorboard --logdir="D:\Twitter Sentiment Analysis\logs\fit"
```

## Index of C:\

Name	Size	Date Modified
\$MfeDeepRem/		2/15/22, 9:03:27 PM
\$Recycle.Bin/		3/16/22, 9:15:18 PM
\$SysReset/		11/30/21, 9:53:17 PM
\$WinREAgent/		6/16/22, 2:28:42 AM
Apps/		12/2/20, 8:00:32 PM
dell/		2/14/21, 12:03:01 PM
Documents and Settings/		12/13/20, 11:44:42 AM
Drivers/		12/2/20, 7:26:26 PM
hp_lj1020_Full_Solution/		3/23/21, 12:06:48 AM
Intel/		6/13/22, 1:40:22 PM
Kruti Dev Fonts(FontBharat.blogspot.com)/		3/22/21, 10:48:22 PM
LJP1100_P1560_P1600_Full_Solution/		1/8/22, 4:47:45 PM
Microsoft Malware Detection/		3/18/22, 4:44:17 AM
MSOCache/		3/26/21, 12:56:21 PM
OneDriveTemp/		2/14/21, 11:41:35 AM
PerfLogs/		6/5/21, 5:40:48 PM
Program Files/		6/12/22, 6:35:49 PM
Program Files (x86)/		2/21/22, 11:05:44 PM
ProgramData/		6/14/22, 11:05:18 PM
Recovery/		10/23/21, 5:51:57 PM
System Volume Information/		5/17/22, 5:10:25 PM
Tomcat/		10/18/21, 10:43:37 PM
Users/		10/23/21, 5:26:51 PM
Windows/		6/18/22, 11:17:17 AM
dell.sdr	28.1 kB	12/2/20, 7:36:02 PM
DumpStack.log	12.0 kB	6/13/22, 1:40:21 PM
DumpStack.log.tmp	12.0 kB	6/18/22, 12:20:17 PM
hiberfil.sys	12.7 GB	6/27/22, 12:13:54 AM
pagefile.sys	48.8 GB	6/18/22, 12:20:17 PM
swapfile.sys	16.0 MB	6/18/22, 12:20:17 PM

**Stacked LSTMs with pre-trained word embeddings**

```
In [135]: def build_LSTM():
            model = tf.keras.models.Sequential()
            model.add(tf.keras.layers.Input(shape=max_length, ))
            model.add(tf.keras.layers.Embedding(len(tokenizer.word_index)+1, EMBEDDING_DIM, weights=[embedding_matrix], input_length=max_length, trainable=True))
            model.add(tf.keras.layers.LSTM(512, return_sequences=True, recurrent_dropout=0.3))
            model.add(tf.keras.layers.LSTM(512, recurrent_dropout=0.3))
            model.add(tf.keras.layers.Dense(len(Y_train.columns), activation='sigmoid'))

            model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])

            return model
```

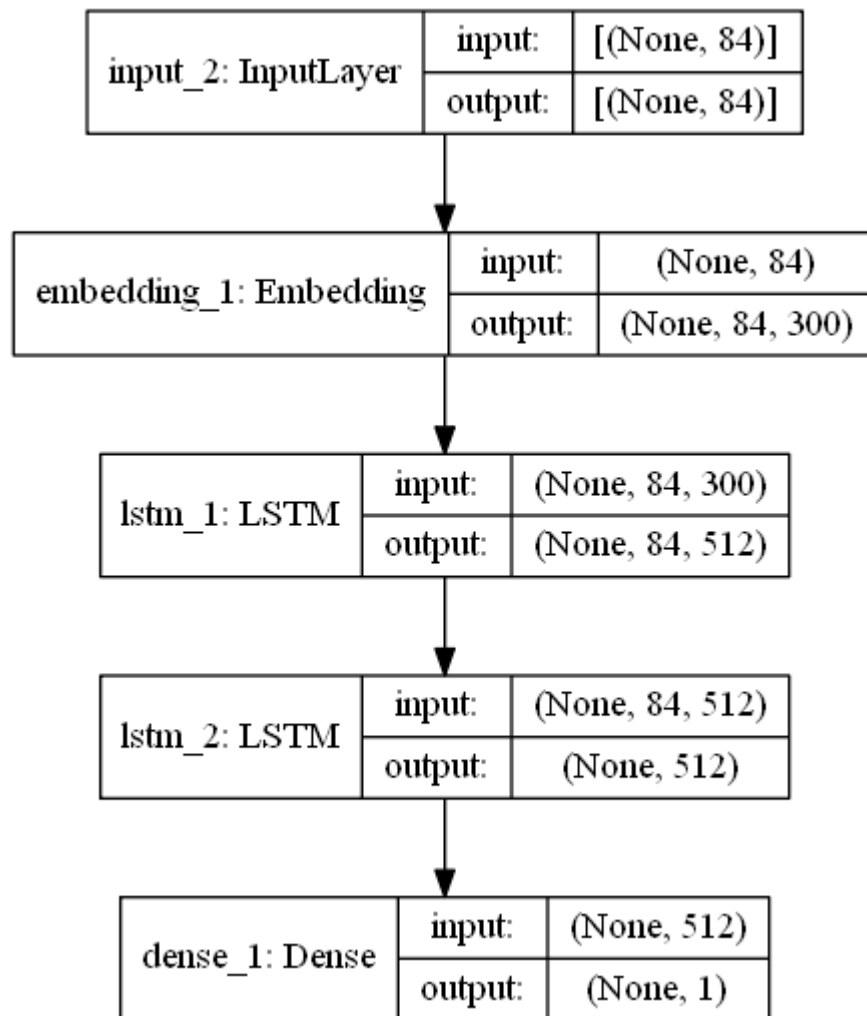
```
In [136]: model_LSTM = build_LSTM()
            model_LSTM.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, 84, 300)	11073000
=====		
lstm_1 (LSTM)	(None, 84, 512)	1665024
=====		
lstm_2 (LSTM)	(None, 512)	2099200
=====		
dense_1 (Dense)	(None, 1)	513
=====		
Total params: 14,837,737		
Trainable params: 14,837,737		
Non-trainable params: 0		
=====		

```
In [137]: tf.keras.utils.plot_model(model_LSTM,to_file='model_LSTM.png', show_shapes=True, show_layer_names=True)
```

Out[137]:



```
In [139]: LSTM_hist = model_LSTM.fit(X_train_pad, Y_train, epochs=10, validation_data=(X_cv_pad, Y_cv), batch_size=64,
                                     callbacks=[earlystop, model_checkpoint_callback, tensorboard_callback])
```

Epoch 1/10

1000/1000 [=====] - 5990s 6s/step - loss: 0.4920 - accuracy: 0.7573 - F1\_score: 0.7286 - Precision: 1.0000 - Recall: 0.5786 - val\_loss: 0.4658 - val\_accuracy: 0.7768 - val\_F1\_score: 0.7355 - val\_Precision: 1.0000 - val\_Recall: 0.5838

Epoch 2/10

1000/1000 [=====] - 5448s 5s/step - loss: 0.4015 - accuracy: 0.8142 - F1\_score: 0.7232 - Precision: 1.0000 - Recall: 0.5703 - val\_loss: 0.4880 - val\_accuracy: 0.7695 - val\_F1\_score: 0.7103 - val\_Precision: 1.0000 - val\_Recall: 0.5532

Epoch 3/10

1000/1000 [=====] - 28814s 29s/step - loss: 0.2892 - accuracy: 0.8715 - F1\_score: 0.7241 - Precision: 1.0000 - Recall: 0.5703 - val\_loss: 0.5642 - val\_accuracy: 0.7554 - val\_F1\_score: 0.7288 - val\_Precision: 1.0000 - val\_Recall: 0.5754

Epoch 00003: early stopping

```
In [140]: loss_lstm, accuracy_lstm, f1_score_lstm, precision_lstm, recall_lstm = model_LSTM.evaluate(X_test_pad, Y_test, verbose=0)
print('F1 score for LSTM model: %f' % f1_score_lstm)

F1 score for LSTM model: 0.720554
```

```
In [141]: # Clear any logs from previous runs
!rm -rf ./logs/

'rm' is not recognized as an internal or external command,
operable program or batch file.
```

In [142]: %tensorboard --logdir="D:\Twitter Sentiment Analysis\logs\fits"

## Index of C:\

Name	Size	Date Modified
\$MfeDeepRem/		2/15/22, 9:03:27 PM
\$Recycle.Bin/		3/16/22, 9:15:18 PM
\$SysReset/		11/30/21, 9:53:17 PM
\$WinREAgent/		6/16/22, 2:28:42 AM
Apps/		12/2/20, 8:00:32 PM
dell/		2/14/21, 12:03:01 PM
Documents and Settings/		12/13/20, 11:44:42 AM
Drivers/		12/2/20, 7:26:26 PM
hp_lj1020_Full_Solution/		3/23/21, 12:06:48 AM
Intel/		6/13/22, 1:40:22 PM
Kruti Dev Fonts(FontBharat.blogspot.com)/		3/22/21, 10:48:22 PM
LJP1100_P1560_P1600_Full_Solution/		1/8/22, 4:47:45 PM
Microsoft Malware Detection/		3/18/22, 4:44:17 AM
MSOCache/		3/26/21, 12:56:21 PM
OneDriveTemp/		2/14/21, 11:41:35 AM
PerfLogs/		6/5/21, 5:40:48 PM
Program Files/		6/12/22, 6:35:49 PM
Program Files (x86)/		2/21/22, 11:05:44 PM
ProgramData/		6/14/22, 11:05:18 PM
Recovery/		10/23/21, 5:51:57 PM
System Volume Information/		5/17/22, 5:10:25 PM
Tomcat/		10/18/21, 10:43:37 PM
Users/		10/23/21, 5:26:51 PM
Windows/		6/18/22, 11:17:17 AM
dell.sdr	28.1 kB	12/2/20, 7:36:02 PM
DumpStack.log	12.0 kB	6/13/22, 1:40:21 PM
DumpStack.log.tmp	12.0 kB	6/18/22, 12:20:17 PM
hiberfil.sys	12.7 GB	6/27/22, 12:13:54 AM
pagefile.sys	48.8 GB	6/18/22, 12:20:17 PM
swapfile.sys	16.0 MB	6/18/22, 12:20:17 PM

**Stacked GRUs with pre-trained word embeddings**

```
In [143]: def build_GRU():
            model = tf.keras.models.Sequential()
            model.add(tf.keras.layers.Input(shape=max_length, ))
            model.add(tf.keras.layers.Embedding(len(tokenizer.word_index)+1, EMBEDDING_DIM, weights=[embedding_matrix], input_length=max_length, trainable=True))
            model.add(tf.keras.layers.GRU(512, return_sequences=True, recurrent_dropout=0.3))
            model.add(tf.keras.layers.GRU(512, recurrent_dropout=0.3))
            model.add(tf.keras.layers.Dense(len(Y_train.columns), activation='sigmoid'))

            model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])

            return model
```

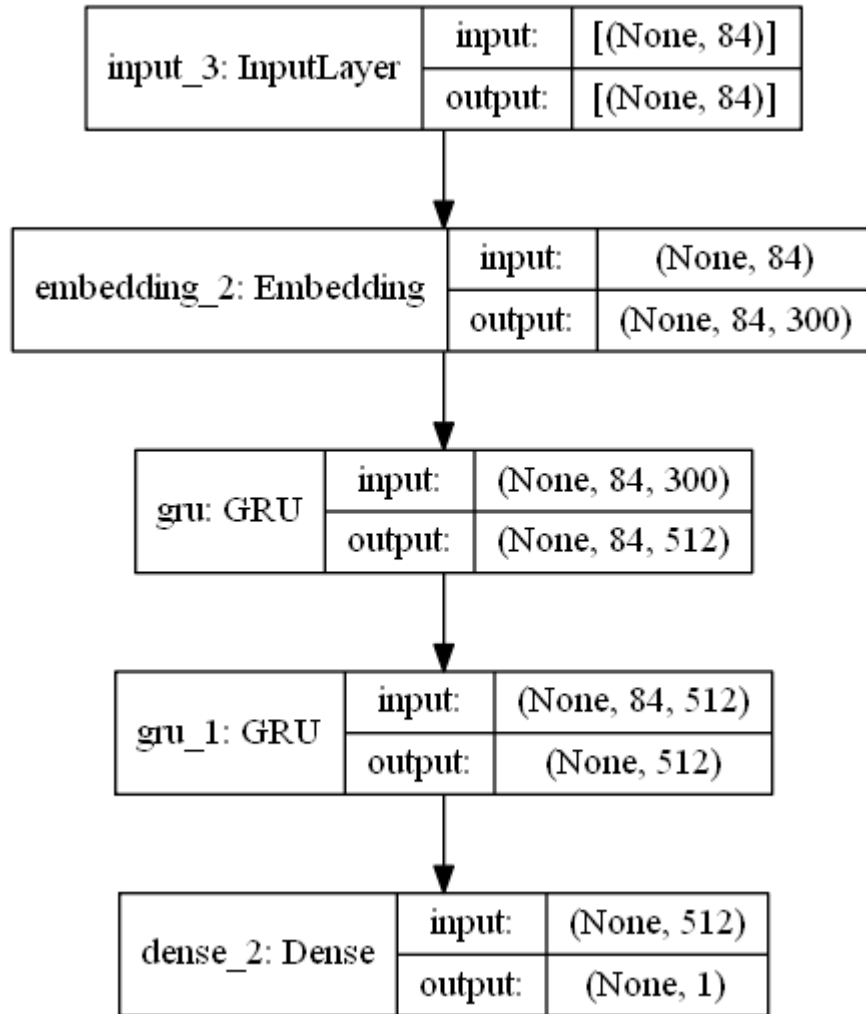
```
In [144]: model_GRU = build_GRU()
            model_GRU.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, 84, 300)	11073000
=====		
gru (GRU)	(None, 84, 512)	1250304
=====		
gru_1 (GRU)	(None, 512)	1575936
=====		
dense_2 (Dense)	(None, 1)	513
=====		
Total params: 13,899,753		
Trainable params: 13,899,753		
Non-trainable params: 0		
=====		

```
In [147]: tf.keras.utils.plot_model(model_GRU,to_file='model_GRU.png', show_shapes=True
, show_layer_names=True)
```

Out[147]:



```
In [149]: GRU_hist = model_GRU.fit(X_train_pad, Y_train, epochs=5, validation_data=(X_cv
_pad, Y_cv), batch_size=64,
                                     callbacks=[earlystop, model_checkpoint_callback, tens
orboard_callback])
```

Epoch 1/5

1000/1000 [=====] - 6655s 7s/step - loss: 0.4923 - accuracy: 0.7576 - F1\_score: 0.7303 - Precision: 0.9990 - Recall: 0.5806 - val\_loss: 0.4672 - val\_accuracy: 0.7739 - val\_F1\_score: 0.7527 - val\_Precision: 1.0000 - val\_Recall: 0.6056

Epoch 2/5

1000/1000 [=====] - 6626s 7s/step - loss: 0.3939 - accuracy: 0.8186 - F1\_score: 0.7223 - Precision: 1.0000 - Recall: 0.5687 - val\_loss: 0.4868 - val\_accuracy: 0.7701 - val\_F1\_score: 0.7400 - val\_Precision: 1.0000 - val\_Recall: 0.5896

Epoch 3/5

1000/1000 [=====] - 7357s 7s/step - loss: 0.2787 - accuracy: 0.8776 - F1\_score: 0.7238 - Precision: 1.0000 - Recall: 0.5700 - val\_loss: 0.5802 - val\_accuracy: 0.7506 - val\_F1\_score: 0.7338 - val\_Precision: 1.0000 - val\_Recall: 0.5815

Epoch 00003: early stopping



```
In [150]: loss_GRU, accuracy_GRU, f1_score_GRU, precision_GRU, recall_GRU = model_GRU.evaluate(X_test_pad, Y_test, verbose=0)
print('F1 score for GRU model: %f' % f1_score_GRU)

F1 score for GRU model: 0.725504
```

## Bidirectional RNN

```
In [156]: def build_biRNN():
    model = tf.keras.models.Sequential()
    model.add(tf.keras.layers.Input(shape=max_length, ))
    model.add(tf.keras.layers.Embedding(len(tokenizer.word_index)+1, EMBEDDING_DIM, weights=[embedding_matrix], input_length=max_length, trainable=True))
    #model.add(tf.keras.layers.Bidirectional(tf.keras.layers.SimpleRNN(512, return_sequences=True, recurrent_dropout=0.3)))
    model.add(tf.keras.layers.Bidirectional(tf.keras.layers.SimpleRNN(512, recurrent_dropout=0.3)))
    model.add(tf.keras.layers.Dense(len(Y_train.columns), activation='sigmoid'))

    model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])

    return model
```

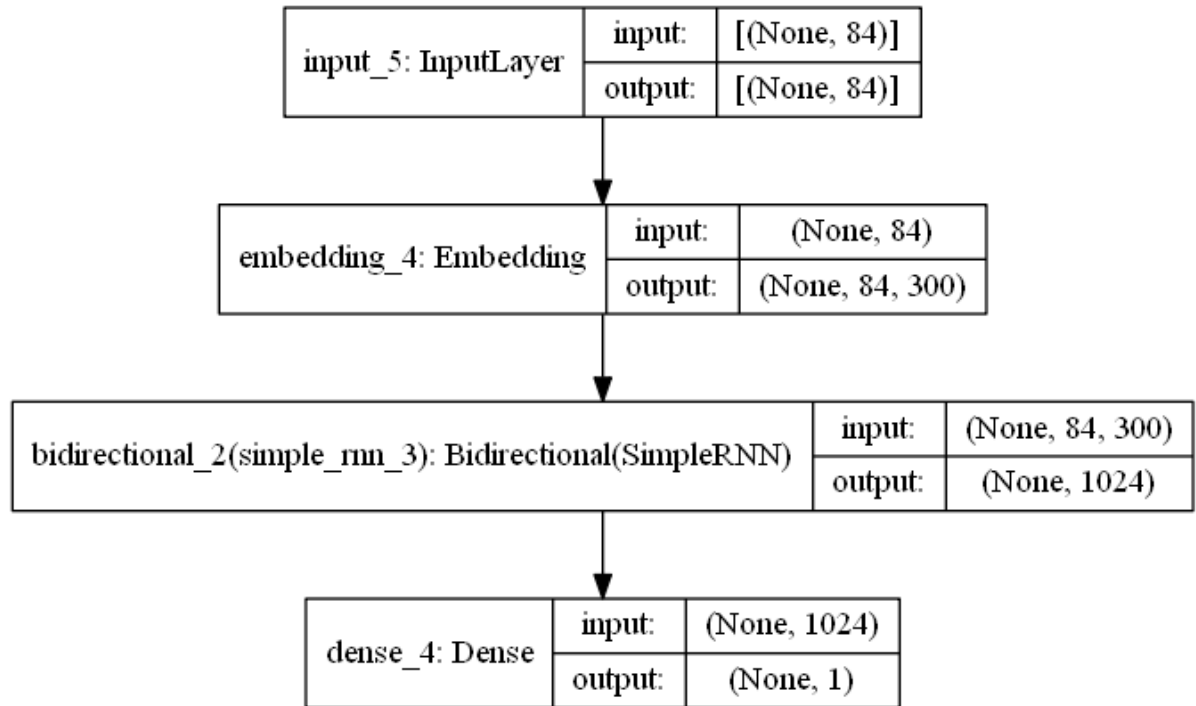
```
In [157]: model_biRNN = build_biRNN()
model_biRNN.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 84, 300)	11073000
bidirectional_2 (Bidirection	(None, 1024)	832512
dense_4 (Dense)	(None, 1)	1025
Total params: 11,906,537		
Trainable params: 11,906,537		
Non-trainable params: 0		

```
In [158]: tf.keras.utils.plot_model(model_biRNN,to_file='model_biRNN.png', show_shapes=True, show_layer_names=True)
```

Out[158]:



```
In [160]: biRNN_hist = model_biRNN.fit(X_train_pad, Y_train, epochs=10, validation_data=(X_cv_pad, Y_cv), batch_size=64, callbacks=[earlystop, model_checkpoint_callback, tensorboard_callback])
```

Epoch 1/10

1000/1000 [=====] - 1982s 2s/step - loss: 0.6101 - accuracy: 0.6713 - F1\_score: 0.7356 - Precision: 1.0000 - Recall: 0.6037 - val\_loss: 0.6180 - val\_accuracy: 0.6779 - val\_F1\_score: 0.4708 - val\_Precision: 1.0000 - val\_Recall: 0.3103

Epoch 2/10

1000/1000 [=====] - 2032s 2s/step - loss: 0.5187 - accuracy: 0.7426 - F1\_score: 0.7306 - Precision: 1.0000 - Recall: 0.5876 - val\_loss: 0.5140 - val\_accuracy: 0.7514 - val\_F1\_score: 0.7595 - val\_Precision: 1.0000 - val\_Recall: 0.6144

Epoch 3/10

1000/1000 [=====] - 1998s 2s/step - loss: 0.5042 - accuracy: 0.7562 - F1\_score: 0.7280 - Precision: 1.0000 - Recall: 0.5830 - val\_loss: 0.5479 - val\_accuracy: 0.7381 - val\_F1\_score: 0.7348 - val\_Precision: 1.0000 - val\_Recall: 0.5830

Epoch 4/10

1000/1000 [=====] - 1651s 2s/step - loss: 0.4291 - accuracy: 0.8052 - F1\_score: 0.7226 - Precision: 1.0000 - Recall: 0.5707 - val\_loss: 0.5413 - val\_accuracy: 0.7470 - val\_F1\_score: 0.7050 - val\_Precision: 1.0000 - val\_Recall: 0.5464

Epoch 00004: early stopping

```
In [161]: loss_biRNN, accuracy_biRNN, f1_score_biRNN, precision_biRNN, recall_biRNN = model_biRNN.evaluate(X_test_pad, Y_test, verbose=0)
print('F1 score for bidirectional RNN model: %f' % f1_score_biRNN)

F1 score for bidirectional RNN model: 0.699530
```

```
In [151]: def build_biLSTM():
    model = tf.keras.models.Sequential()
    model.add(tf.keras.layers.Input(shape=max_length, ))
    model.add(tf.keras.layers.Embedding(len(tokenizer.word_index)+1, EMBEDDING_DIM, weights=[embedding_matrix], input_length=max_length, trainable=True))
    model.add(tf.keras.layers.Bidirectional(tf.keras.layers.SimpleRNN(512, return_sequences=True, recurrent_dropout=0.3)))
    model.add(tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(512, recurrent_dropout=0.3)))
    model.add(tf.keras.layers.Dense(len(Y_train.columns), activation='sigmoid'))

    model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])

    return model
```

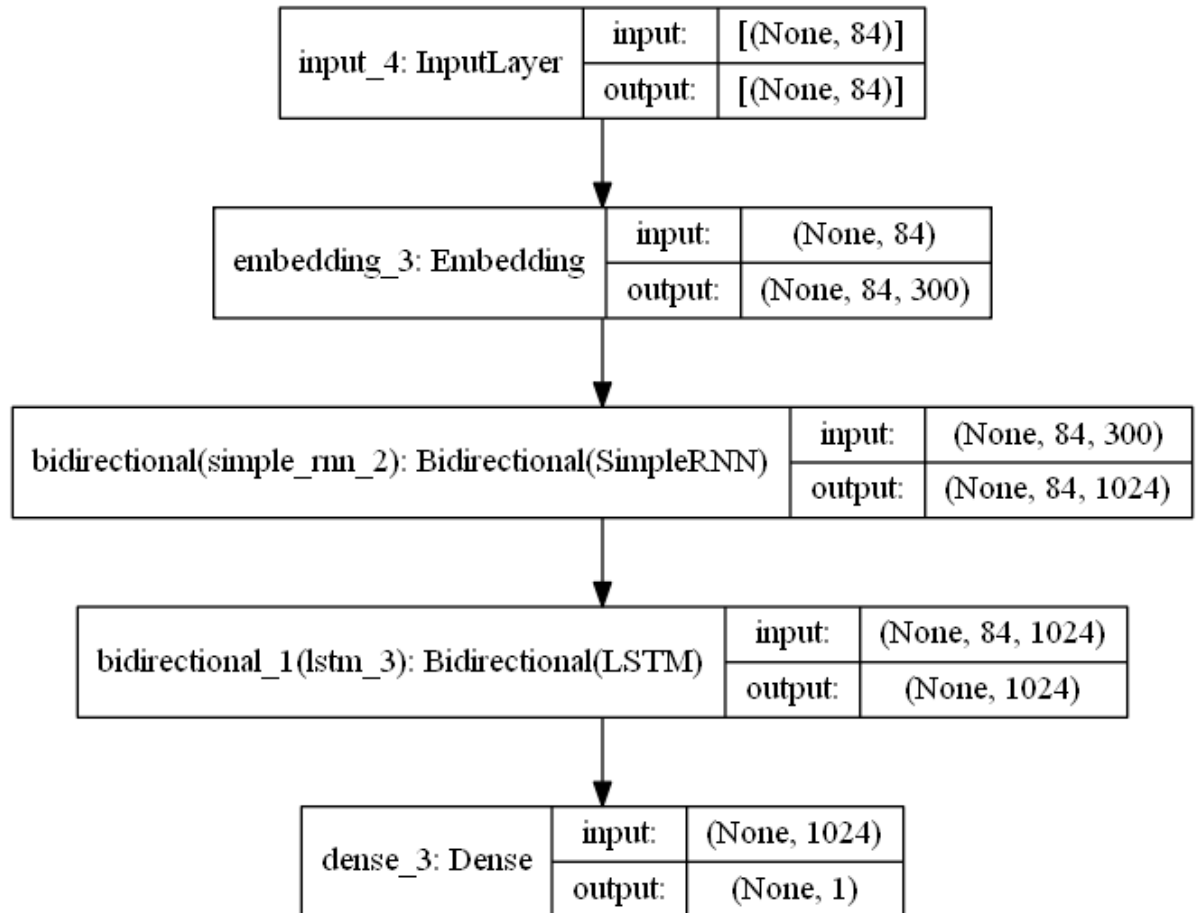
```
In [152]: model_biLSTM = build_biLSTM()
model_biLSTM.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, 84, 300)	11073000
-----		
bidirectional (Bidirectional)	(None, 84, 1024)	832512
-----		
bidirectional_1 (Bidirectional)	(None, 1024)	6295552
-----		
dense_3 (Dense)	(None, 1)	1025
=====		
Total params: 18,202,089		
Trainable params: 18,202,089		
Non-trainable params: 0		
-----		

```
In [153]: tf.keras.utils.plot_model(model_biLSTM,to_file='model_biLSTM.png', show_shapes
= True, show_layer_names=True)
```

Out[153]:



```
In [154]: biRLSTM_hist = model_biLSTM.fit(X_train_pad, Y_train, epochs=10, validation_data=(X_cv_pad, Y_cv), batch_size=64,
                                           callbacks=[earlystop, model_checkpoint_callback, tensorboard_callback])
```

```
Epoch 1/10
1000/1000 [=====] - 15409s 15s/step - loss: 0.5437 - accuracy: 0.7210 - F1_score: 0.7272 - Precision: 0.9990 - Recall: 0.5834 - val_loss: 0.5169 - val_accuracy: 0.7514 - val_F1_score: 0.6405 - val_Precision: 1.0000 - val_Recall: 0.4735
Epoch 2/10
1000/1000 [=====] - 45553s 46s/step - loss: 0.4479 - accuracy: 0.7886 - F1_score: 0.7184 - Precision: 1.0000 - Recall: 0.5656 - val_loss: 0.4987 - val_accuracy: 0.7663 - val_F1_score: 0.7372 - val_Precision: 1.0000 - val_Recall: 0.5859
Epoch 3/10
1000/1000 [=====] - 16117s 16s/step - loss: 0.3895 - accuracy: 0.8217 - F1_score: 0.7173 - Precision: 1.0000 - Recall: 0.5630 - val_loss: 0.5203 - val_accuracy: 0.7574 - val_F1_score: 0.7510 - val_Precision: 1.0000 - val_Recall: 0.6035
Epoch 4/10
1000/1000 [=====] - 15619s 16s/step - loss: 0.3406 - accuracy: 0.8473 - F1_score: 0.7184 - Precision: 1.0000 - Recall: 0.5640 - val_loss: 0.5644 - val_accuracy: 0.7489 - val_F1_score: 0.7017 - val_Precision: 1.0000 - val_Recall: 0.5427
Epoch 00004: early stopping
```

```
In [155]: loss_biLSTM, accuracy_biLSTM, f1_score_biLSTM, precision_biLSTM, recall_biLSTM = model_biLSTM.evaluate(X_test_pad, Y_test, verbose=0)
print('F1 score for bidirectional LSTM model: %f' % f1_score_biLSTM)
```

```
F1 score for bidirectional LSTM model: 0.696684
```

## BERT

```
In [67]: bert_preprocessor = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/1")
bert_encoder = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/3", trainable=True)
```

```
In [68]: """bert_preprocess = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3")
bert_encoder = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/4")"""
```

```
Out[68]: 'bert_preprocess = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3")\nbert_encoder = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/4")'
```

```
In [94]: # Bert Layers
text_input = tf.keras.layers.Input(shape=(), dtype=tf.string, name='text')
preprocessed_text = bert_preprocessor(text_input)
outputs = bert_encoder(preprocessed_text)
pooled_output = outputs["pooled_output"] # [batch_size, 768].
sequence_output = outputs["sequence_output"] # [batch_size, seq_length, 768].

# Neural network layers
#l = tf.keras.layers.Dropout(0.3, name="dropout")(outputs['pooled_output'])
l = tf.keras.layers.Dropout(0.3, name="dropout")(pooled_output)
l = tf.keras.layers.Dense(1, activation='sigmoid', name="output")(l)

# Use inputs and outputs to construct a final model
model_BERT = tf.keras.Model(inputs=[text_input], outputs = [l])
```

```
In [95]: model_BERT.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
text (InputLayer)	[(None,)]	0	
keras_layer (KerasLayer)	{'input_type_ids': ( 0		text[0][0]
keras_layer_1 (KerasLayer)	{'encoder_outputs': 109482241		keras_layer
[2][0]			keras_layer
[2][1]			keras_layer
[2][2]			
dropout (Dropout)	(None, 768)	0	keras_layer_1[2][13]
output (Dense)	(None, 1)	769	dropout[0]
Total params: 109,483,010			
Trainable params: 109,483,009			
Non-trainable params: 1			

```
In [96]: model_BERT.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy', F1_score, Precision, Recall])
```

```
In [100]: BERT_hist = model_BERT.fit(X_train, Y_train, epochs=5, validation_data=(X_cv, Y_cv), batch_size=64,
                                     , callbacks=[earlystop, model_checkpoint_callback, tensorboard_callback])
```

```
Epoch 1/5
1000/1000 [=====] - 49372s 49s/step - loss: 0.7161 - accuracy: 0.5361 - F1_score: 0.8445 - Precision: 0.9950 - Recall: 0.7800 - val_loss: 0.6851 - val_accuracy: 0.5646 - val_F1_score: 1.0000 - val_Precision: 1.0000 - val_Recall: 1.0000
Epoch 2/5
1000/1000 [=====] - 69878s 70s/step - loss: 0.6889 - accuracy: 0.5538 - F1_score: 0.9471 - Precision: 1.0000 - Recall: 0.9217 - val_loss: 0.6849 - val_accuracy: 0.5646 - val_F1_score: 1.0000 - val_Precision: 1.0000 - val_Recall: 1.0000
Epoch 3/5
1000/1000 [=====] - 64981s 65s/step - loss: 0.6868 - accuracy: 0.5613 - F1_score: 0.9832 - Precision: 0.9990 - Recall: 0.9750 - val_loss: 0.6849 - val_accuracy: 0.5646 - val_F1_score: 1.0000 - val_Precision: 1.0000 - val_Recall: 1.0000
Epoch 4/5
1000/1000 [=====] - 66188s 66s/step - loss: 0.6860 - accuracy: 0.5645 - F1_score: 0.9996 - Precision: 1.0000 - Recall: 0.9992 - val_loss: 0.6852 - val_accuracy: 0.5646 - val_F1_score: 1.0000 - val_Precision: 1.0000 - val_Recall: 1.0000
Epoch 5/5
1000/1000 [=====] - 50334s 50s/step - loss: 0.6859 - accuracy: 0.5636 - F1_score: 0.9942 - Precision: 1.0000 - Recall: 0.9925 - val_loss: 0.6858 - val_accuracy: 0.5646 - val_F1_score: 1.0000 - val_Precision: 1.0000 - val_Recall: 1.0000
Epoch 00005: early stopping
```

```
In [102]: loss_BERT, accuracy_BERT, f1_score_BERT, precision_BERT, recall_BERT = model_BERT.evaluate(X_test, Y_test, verbose=0)
          print('F1 score for BERT model: %f' % f1_score_BERT)
```

```
F1 score for BERT model: 1.000000
```

## Summary

```
In [166]: # assigning data
table_data = [("Logistic Regression", f1_score_log, "accuracy_lg")
              ,("Naive Bayes", f1_score_NB, "accuracy_NB")
              ,("Random Forest", f1_score_RF, "accuracy_RF")
              ,("Gradient Boosting Decision Trees", f1_score_GBDT, "accuracy_G
BDT")
              ,("Linear SVC", f1_score_SVC, "accuracy_SVC")
              ,("RNN", f1_score_rnn, accuracy_rnn)
              ,("LSTM", f1_score_lstm, accuracy_lstm)
              ,("GRU", f1_score_GRU, accuracy_GRU)
              ,("Bidirectional RNN", f1_score_biRNN, accuracy_biRNN)
              ,("Bidirectional LSTM", f1_score_biLSTM, accuracy_biLSTM)
              ,("BERT", f1_score_BERT, accuracy_BERT, accuracy_BERT)
              ]

# creating header
table_head = ["Model", "F1 Score", "Accuracy"]

# displaying the table
print(tabulate(table_data, headers=table_head, tablefmt="grid"))
```

Model	F1 Score	Accuracy
Logistic Regression	0.786995	accuracy_lg
Naive Bayes	0.784441	accuracy_NB
Random Forest	0.777063	accuracy_RF
Gradient Boosting Decision Trees	0.787067	accuracy_GBDT
Linear SVC	0.788652	accuracy_SVC
RNN	0.513728	0.702170193195343
LSTM	0.720554	0.7546254396438599
GRU	0.725504	0.7545254230499268
Bidirectional RNN	0.69953	0.7496249675750732
Bidirectional LSTM	0.696684	0.7431243062019348
BERT	1	0.5646564364433289

Conclusion: As we can see from the above table, the f1 score of BERT is the highest but it does not have the highest accuracy.